




LoRA: 5 adapters $\times$ 30 blocks = 2,334,720 trainable parameters, 0.18% of the 1.3B texture flow. Temporal attention adds none.